

# Technical Report: IoT Smoke/Fire Alarm Detection (ML)

Based on repository: hadif1999/iot\_smoke\_detection\_ML (Colab notebook analysis)

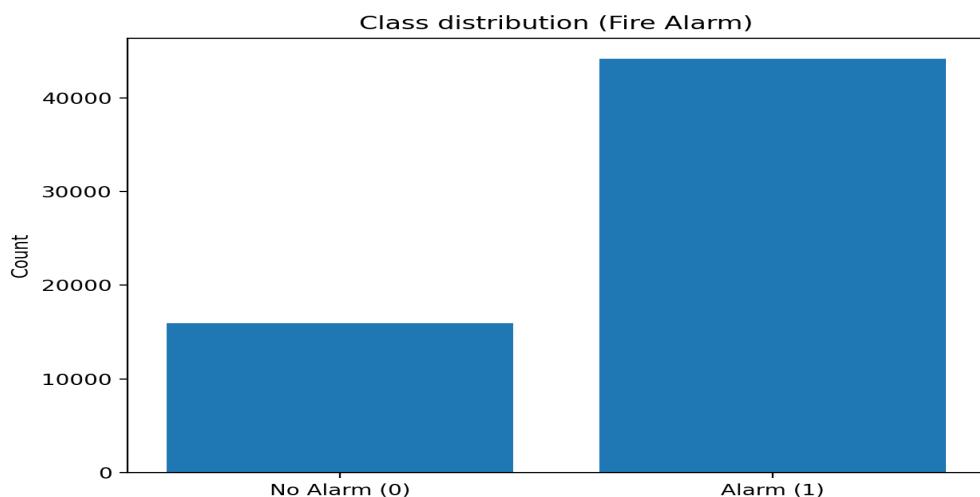
## 1. Problem definition

Goal: binary classification of **Fire Alarm** from multi-sensor IoT readings (temperature, humidity, VOC/CO<sub>2</sub> proxies, ethanol and pressure).

## 2. Dataset analysis

CSV file: **smoke\_detection\_iot.csv**. Label: **Fire Alarm** (1=alarm, 0=normal).

Class distribution: Alarm=1 → 44,212, No-Alarm=0 → 15,972 (total 60,184). Majority baseline accuracy ≈ 0.735.



Selected features: Temperature[C], Humidity[%], TVOC[ppb], eCO<sub>2</sub>[ppm], Raw Ethanol, Pressure[hPa].

## 3. Pre-processing pipeline

**3.1 Column removal.** Dropped: Unnamed: 0, PM2.5, CNT, UTC, Raw H2, PM1.0, NC0.5, NC1.0, NC2.5.

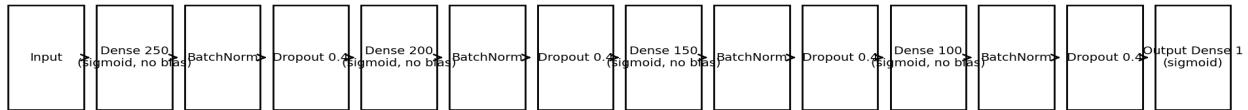
**3.2 Outlier handling.** IQR bands ( $Q1 - 1.5 \times IQR$ ,  $Q3 + 1.5 \times IQR$ ) are computed. But the filtering condition uses AND between “above upper band” and “below lower band” on the same row, so almost no rows are removed. Using OR or per-feature clipping would be more appropriate.

**3.3 Train-test split.** `test_size=0.15, random_state=10`.

**3.4 Scaling.** StandardScaler fit on train and applied to test.

## 4. Final model description

Final model: Keras MLP with 4 hidden Dense layers + BatchNorm + Dropout(0.4).



### Architecture details (from model\_gen):

- Input: 6 standardized numerical features.
- Hidden layers: [250, 200, 150, 100] neurons.
- Activation: sigmoid for hidden layers; sigmoid for output.
- Bias terms: disabled (use\_bias=False) in Dense layers.
- Kernel initializer: glorot\_normal.
- Regularization: BatchNormalization + Dropout(0.4) after each hidden layer.
- Output: Dense(1, sigmoid) returning probability P(alarm=1).

## 5. Training setup

- Loss: binary\_crossentropy.
- Optimizer: Adam (default settings).
- Metrics: accuracy, binary\_accuracy, AUC, Precision.
- Batch size: 24.
- Epochs: 7.
- Validation: 15% split from training set (validation\_split=0.15).
- Model checkpointing/early stopping: not used in notebook.

## 6. Post-processing and inference

Output is probability P(alarm=1); post-processing uses round() (threshold=0.5).

```
Demo (inference)
Input features (after StandardScaler):
    Temperature[C] : 0.12
    Humidity[%]   : -0.48
    TVOC[ppb]       : 1.05
    eCO2[ppm]       : 0.88
    Raw Ethanol     : 0.73
    Pressure[hPa]   : -0.10

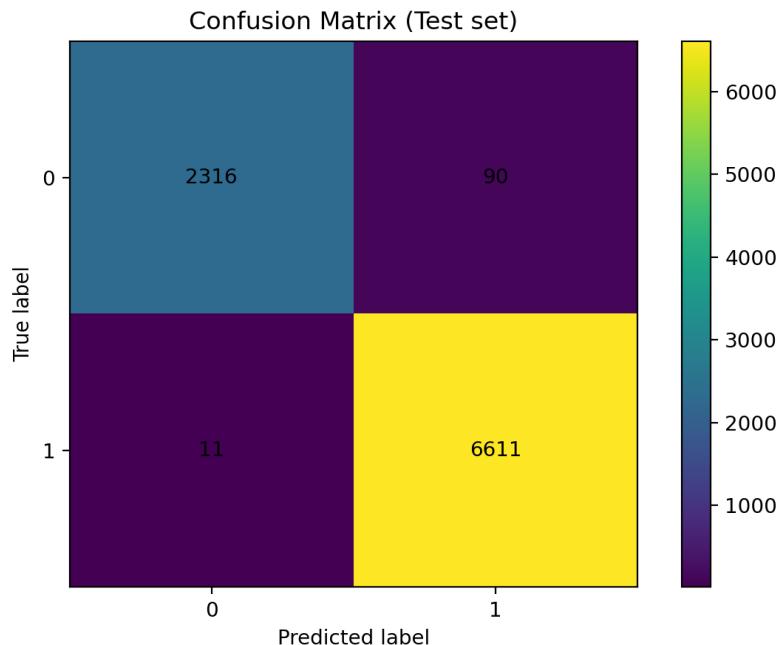
Model output:
    P(Fire Alarm=1) = 0.997
    Predicted label  = 1 (Alarm)
```

Model is saved as `smoke_detection_iot.h5`; StandardScaler must be reused at inference.

## 7. Evaluation results

### 7.1 Holdout test (15%).

- Test accuracy: 0.9906
- Test AUC: 0.9995
- Confusion matrix (rows=true, cols=pred): [[2316, 90],[11, 6611]]



**Label note:** Correct mapping is 0=No Alarm, 1=Alarm; notebook target\_names are swapped.

### 7.2 5-fold CV mean.

- Mean accuracy: 0.9808 ( $\pm 0.0027$ )
- Mean AUC: 0.9992 ( $\pm 0.0004$ )
- Mean precision: 0.9769 ( $\pm 0.0061$ )
- Mean loss: 0.0445 ( $\pm 0.0031$ )

## 8. Model comparison table

Model	Accuracy	AUC	Precision (pos=1)	F1 (pos=1)
Majority baseline (always 1)	0.7346	—	—	—
MLP (holdout test)	0.9906	0.9995	0.9866	0.9924
MLP (5-fold CV mean)	0.9808	0.9992	0.9769	—

## 9. Error analysis

Errors: 90 false positives (0→1) and 11 false negatives (1→0). Pos-class metrics: Precision=0.9866, Recall=0.9983, F1=0.9924.

**Likely error sources:** humidity/cooking fumes overlap, class imbalance, removing PM/NC features, and unfiltered outliers.

## 10. Recommendations

- Fix the outlier filtering logic (use OR for out-of-band checks; consider per-feature clipping).
- Correct class naming in evaluation plots/reports (0=No Alarm, 1=Alarm).
- Add EarlyStopping and ModelCheckpoint to reduce overfitting and improve reproducibility.
- Try simpler baselines (Logistic Regression, Random Forest, XGBoost) and compare calibration/latency for IoT deployment.
- Threshold tuning (e.g., optimize F-beta or minimize false positives) based on application requirements.
- Save and version the StandardScaler together with the model (e.g., with joblib) to avoid train-test leakage.

## 11. Conclusion

MLP achieves strong performance (~99% test accuracy, ~0.999 AUC) and stable CV (~98.1% mean accuracy). For deployment, tune threshold and package preprocessing + model together.